

Homework 4

Chen Yi-Ju(Ernie)

2020/6/9

Question 9.1

Question 9.1

Using the same crime data set `uscrime.txt` as in Question 8.2, apply Principal Component Analysis and then create a regression model using the first few principal components. Specify your new model in terms of the original variables (not the principal components), and compare its quality to that of your solution to Question 8.2.

Answer: Using an estimation using both 4 and 5 PCAs (assumption based on scree plot), the results were worse than the original regression. This may be due to overtraining on the original regression.

```
setwd("D:/ernie/self-study/GTxMicroMasters/Introduction to Analytics Modeling/week4/homework")
library(tidyverse)
crime <- read.table("uscrime.txt",header = T) %>%
  data.frame()

#original model from 8.2 with 9 variables
model <- lm (data = crime , Crime ~ Ed + Ineq + LF + M + M.F + Po1 + Pop + Prob + Time )
summary(model)
```

```
##
## Call:
## lm(formula = Crime ~ Ed + Ineq + LF + M + M.F + Po1 + Pop + Prob +
##      Time, data = crime)
##
## Residuals:
##      Min       1Q   Median       3Q      Max
## -468.62 -100.73   -6.44   139.91   520.35
##
## Coefficients:
##              Estimate Std. Error t value Pr(>|t|)
## (Intercept) -5189.5782   1460.9341  -3.552 0.001063 **
## Ed           140.9730     57.8900    2.435 0.019823 *
## Ineq         68.7477     15.8765    4.330 0.000109 ***
## LF          -609.2340   1065.0117  -0.572 0.570751
## M            68.3357     35.3331    1.934 0.060784 .
## M.F          17.8666     15.2790    1.169 0.249738
## Po1          126.5215     17.3893    7.276 1.22e-08 ***
## Pop          -0.6526     1.2716   -0.513 0.610833
## Prob        -4006.6838   2033.8562  -1.970 0.056359 .
## Time           1.7858      6.6248    0.270 0.788995
```

```
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
##
## Residual standard error: 213.8 on 37 degrees of freedom
## Multiple R-squared:  0.7543, Adjusted R-squared:  0.6945
## F-statistic: 12.62 on 9 and 37 DF,  p-value: 7.275e-09

pca <- prcomp(formula = ~ So + Ed + Ineq + LF + M + M.F + Po1 + Pop + Prob + Time + Po2 + NW + U1 + U2 + Wealth, data = data)
pca

## Standard deviations (1, ..., p=15):
## [1] 2.45335539 1.67387187 1.41596057 1.07805742 0.97892746 0.74377006
## [7] 0.56729065 0.55443780 0.48492813 0.44708045 0.41914843 0.35803646
## [13] 0.26332811 0.24180109 0.06792764
##
## Rotation (n x k) = (15 x 15):
##
```

	PC1	PC2	PC3	PC4	PC5
## So	-0.33088129	0.15837219	-0.0155433104	-0.29247181	0.12061130
## Ed	0.33962148	-0.21461152	-0.0677396249	-0.07974375	0.02442839
## Ineq	-0.36579778	0.02752240	0.0002944563	0.08066612	0.21672823
## LF	0.17617757	-0.31943042	-0.2715301768	0.14326529	0.39407588
## M	-0.30371194	-0.06280357	-0.1724199946	0.02035537	0.35832737
## M.F	0.11638221	-0.39434428	0.2031621598	-0.01048029	0.57877443
## Po1	0.30863412	0.26981761	-0.0506458161	-0.33325059	0.23527680
## Pop	0.11307836	0.46723456	-0.0770210971	0.03210513	0.08317034
## Prob	-0.25888661	-0.15831708	0.1176726436	-0.49303389	-0.16562829
## Time	-0.02062867	0.38014836	-0.2235664632	0.54059002	0.14764767
## Po2	0.31099285	0.26396300	-0.0530651173	-0.35192809	0.20473383
## NW	-0.29358647	0.22801119	-0.0788156621	-0.23925971	0.36079387
## U1	0.04050137	-0.00807439	0.6590290980	0.18279096	0.13136873
## U2	0.01812228	0.27971336	0.5785006293	0.06889312	0.13499487
## Wealth	0.37970331	0.07718862	-0.0100647664	-0.11781752	-0.01167683

```
##
```

	PC6	PC7	PC8	PC9	PC10	PC11
## So	-0.100500743	0.19649727	-0.22734157	0.65599872	-0.06141452	-0.23397868
## Ed	-0.008571367	-0.23943629	0.14644678	0.44326978	-0.51887452	0.11821954
## Ineq	0.272027031	0.37483032	-0.07184018	0.02494384	0.01390576	0.18278697
## LF	0.504234275	-0.15931612	-0.25513777	-0.14393498	-0.03077073	-0.38532827
## M	-0.449132706	-0.15707378	0.55367691	-0.15474793	0.01443093	-0.39446657
## M.F	-0.074501901	0.15548197	0.05507254	0.24378252	0.35323357	0.28029732
## Po1	-0.095776709	0.08011735	-0.04613156	-0.19425472	0.14320978	0.13042001
## Pop	0.547098563	0.09046187	0.59078221	0.20244830	0.03970718	-0.05849643
## Prob	0.283535996	-0.56159383	0.08598908	0.05306898	0.42530006	0.08978385
## Time	-0.148203050	-0.44199877	-0.19507812	0.23551363	0.29264326	0.26363121
## Po2	-0.119524780	0.09518288	-0.03168720	-0.19512072	0.05929780	0.13885912
## NW	0.051219538	-0.31154195	-0.20432828	-0.18984178	-0.49201966	0.20695666
## U1	0.017385981	-0.17354115	0.20206312	-0.02069349	-0.22765278	0.17857891
## U2	0.048155286	-0.07526787	-0.24369650	-0.05576010	0.04750100	-0.47021842
## Wealth	-0.154683104	-0.14859424	-0.08630649	0.23196695	0.11219383	-0.31955631


```
##
```

	PC12	PC13	PC14	PC15
## So	0.05753357	-0.29368483	0.29364512	0.0084369230
## Ed	-0.47786536	0.19441949	-0.03964277	-0.0280052040
## Ineq	-0.43762828	-0.12181090	-0.59279037	0.0177570357
## LF	-0.02705134	-0.27742957	0.15385625	0.0336823193
## M	-0.16580189	-0.05142365	-0.04901705	0.0051398012
## M.F	0.23925913	0.31624667	0.04125321	0.0097922075

```
## Po1      -0.22611207 -0.18592255  0.09490151 -0.6894155129
## Pop       0.18350385  0.12651689  0.05326383  0.0001496323
## Prob     -0.15567100 -0.03547596 -0.04761011  0.0293376260
## Time     -0.13536989 -0.05738113  0.04488401  0.0376754405
## Po2     -0.19088461 -0.13454940  0.08259642  0.7200270100
## NW       0.36671707  0.22901695 -0.13227774 -0.0370783671
## U1       0.09314897 -0.59039450  0.02335942  0.0111359325
## U2      -0.28440496  0.43292853  0.03985736  0.0073618948
## Wealth   0.32172821 -0.14077972 -0.70031840 -0.0025685109
```

```
#Plotting Scree plot
# Kaiser eigenvalue-greater-than-one rule
Scree <- plot(pca,
              type="line",
              main="Scree Plot for crime factors")%>%
  abline(h=1, col="blue")
```

Scree Plot for crime factors




```
# Calculate and Plot the variances and proportion of variances
```


```
var <- pca$sdev^2
propvar <- var/sum(var)
```

```
#plotting
```

```
por.explained <- plot(propvar, xlab = "Principle component" , ylab = "Porpotion explained" , type = "s")
```



```
acc.por.explained <- plot(cumsum(propvar) , xlab = "Principle component" , ylab = "AccamulatedPorportion explained")
```



```
#choose 4 new variables(variance > 1)
pca.chosen <- pca$x[, 1:4]
pca.chosen
```

##	PC1	PC2	PC3	PC4
## 1	-4.1992835	1.09383120	1.11907395	-0.67178115
## 2	1.1726630	-0.67701360	0.05244634	0.08350709
## 3	-4.1737248	-0.27677501	0.37107658	-0.37793995
## 4	3.8349617	2.57690596	-0.22793998	-0.38262331
## 5	1.8392999	-1.33098564	-1.27882805	-0.71814305
## 6	2.9072336	0.33054213	-0.53288181	-1.22140635
## 7	0.2457752	0.07362562	0.90742064	-1.13685873
## 8	-0.1301330	1.35985577	-0.59753132	-1.44045387
## 9	-3.6103169	0.68621008	-1.28372246	-0.55171150
## 10	1.1672376	-3.03207033	-0.37984502	0.28887026
## 11	2.5384879	2.66771358	-1.54424656	0.87671210
## 12	1.0065920	0.06044849	-1.18861346	1.31261964
## 13	0.5161143	-0.97485189	-1.83351610	1.59117618
## 14	0.4265556	-1.85044812	-1.02893477	0.07789173
## 15	-3.3435299	-0.05182823	1.01358113	-0.08840211
## 16	-3.0310689	2.10295524	1.82993161	-0.52347187
## 17	-0.2262961	-1.44939774	1.37565975	-0.28960865
## 18	-0.1127499	0.39407030	0.38836278	-3.97985093
## 19	2.9195668	1.58646124	-0.97612613	-0.78629766
## 20	2.2998485	1.73396487	2.82423222	0.23281758
## 21	1.1501667	-0.13531015	-0.28506743	2.19770548

```
## 22 -5.6594827 1.09730404 -0.10043541 0.05245484
## 23 -0.1011749 0.57911362 -0.71128354 0.44394773
## 24 1.3836281 -1.95052341 2.98485490 0.35942784
## 25 0.2727756 -2.63013778 -1.83189535 -0.05207518
## 26 4.0565577 -1.17534729 0.81690756 -1.66990720
## 27 0.8929694 -0.79236692 -1.26822542 0.57575615
## 28 0.1514495 -1.44873320 -0.10857670 0.51040146
## 29 3.5592481 4.76202163 -0.75080576 -0.64692974
## 30 -4.1184576 0.38073981 -1.43463965 -0.63330834
## 31 -0.6811731 -1.66926027 2.88645794 1.30977099
## 32 1.7157269 1.30836339 0.55971313 0.70557980
## 33 -1.8860627 -0.59058174 -1.43570145 -0.18239089
## 34 1.9526349 -0.52395429 0.75642216 -0.44289927
## 35 1.5888864 3.12998571 1.73107199 1.68604766
## 36 1.0709414 1.65628271 -0.79436888 1.85172698
## 37 -4.1101715 -0.15766712 -2.36296974 0.56868399
## 38 -0.7254706 -2.89263339 0.36348376 0.50612576
## 39 -3.3451254 0.95045293 -0.19551398 0.27716645
## 40 -1.0644466 1.05265304 -0.82886286 0.12042931
## 41 1.4933989 -1.86712106 -1.81853582 1.06112429
## 42 -0.6789284 -1.83156328 1.65435992 -0.95121379
## 43 -2.4164258 0.46701087 -1.42808323 -0.41149015
## 44 2.2978729 -0.41865689 0.64422929 0.63462770
## 45 -2.9245282 1.19488555 3.35139309 1.48966984
## 46 1.7654525 -0.95655926 -0.98576138 -1.05683769
## 47 2.3125056 -2.56161119 1.58223354 -0.59863946
```

#Combining PCAs with crime

```
new.crime <- cbind(pca.chosen, crime[,16]) %>%
  data.frame()
colnames(new.crime)[5] <- "Crime"
new.model <- lm ( Crime ~ PC1 + PC2 + PC3 + PC4 ,data = new.crime )
summary(new.model)
```

```
##
## Call:
## lm(formula = Crime ~ PC1 + PC2 + PC3 + PC4, data = new.crime)
##
## Residuals:
##      Min       1Q   Median       3Q      Max
## -557.76 -210.91  -29.08   197.26   810.35
##
## Coefficients:
##              Estimate Std. Error t value Pr(>|t|)
## (Intercept)    905.09      49.07   18.443 < 2e-16 ***
## PC1              65.22      20.22    3.225  0.00244 **
## PC2              70.08      29.63    2.365  0.02273 *
## PC3             -25.19      35.03   -0.719  0.47602
## PC4             -69.45      46.01   -1.509  0.13872
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
##
## Residual standard error: 336.4 on 42 degrees of freedom
## Multiple R-squared:  0.3091, Adjusted R-squared:  0.2433
## F-statistic: 4.698 on 4 and 42 DF,  p-value: 0.003178
```

```
#Transforming new data back to original variable
```

```
PCAs <- new.model$coefficients[2:5]  
intercept <- new.model$coefficients[1]  
original <- pca$rotation[,1:4] %*% PCAs  
original
```

```
##           [,1]  
## So      10.223091  
## Ed      14.352610  
## Ineq   -27.536348  
## LF     -14.005349  
## M      -21.277963  
## M.F    -24.437572  
## Po1     63.456426  
## Pop     39.830667  
## Prob     3.295707  
## Time    -6.612616  
## Po2     64.557974  
## NW      15.434545  
## U1     -27.222281  
## U2       1.425902  
## Wealth  38.607855
```

```
PCAs
```

```
##      PC1      PC2      PC3      PC4  
## 65.21593 70.08312 -25.19408 -69.44603
```

These are the variables used, expressed in original form.

```
# un-scaling data
```

```
origi.var <- original/sapply(crime[,1:15],sd)  
origi.inter <- intercept - sum(original*sapply(crime[,1:15],mean)/sapply(crime[,1:15],sd))  
origi.var
```

```
##           [,1]  
## So      8.13445978  
## Ed     29.96524916  
## Ineq   -24.61459872  
## LF     -4.71259516  
## M      -7.60978529  
## M.F   -604.71355665  
## Po1     21.53447549  
## Pop     1.04621551  
## Prob     0.32050426  
## Time  -366.78104113  
## Po2     76.44113075  
## NW      0.01599585  
## U1     -6.82330056  
## U2     62.71293220  
## Wealth  5.44778133
```

```
#Trying with 5 PCAs
#choose 4 new variables(variance > 1)
pca.chosen2 <- pca$x[, 1:5]
pca.chosen2
```

##	PC1	PC2	PC3	PC4	PC5
## 1	-4.1992835	1.09383120	1.11907395	-0.67178115	-0.055283376
## 2	1.1726630	-0.67701360	0.05244634	0.08350709	1.173199821
## 3	-4.1737248	-0.27677501	0.37107658	-0.37793995	-0.541345246
## 4	3.8349617	2.57690596	-0.22793998	-0.38262331	1.644746496
## 5	1.8392999	-1.33098564	-1.27882805	-0.71814305	-0.041590320
## 6	2.9072336	0.33054213	-0.53288181	-1.22140635	-1.374360960
## 7	0.2457752	0.07362562	0.90742064	-1.13685873	-0.718644387
## 8	-0.1301330	1.35985577	-0.59753132	-1.44045387	0.222781388
## 9	-3.6103169	0.68621008	-1.28372246	-0.55171150	0.324292990
## 10	1.1672376	-3.03207033	-0.37984502	0.28887026	0.646056610
## 11	2.5384879	2.66771358	-1.54424656	0.87671210	0.324083561
## 12	1.0065920	0.06044849	-1.18861346	1.31261964	-0.358087724
## 13	0.5161143	-0.97485189	-1.83351610	1.59117618	-0.599881946
## 14	0.4265556	-1.85044812	-1.02893477	0.07789173	-0.741887592
## 15	-3.3435299	-0.05182823	1.01358113	-0.08840211	-0.002969448
## 16	-3.0310689	2.10295524	1.82993161	-0.52347187	0.387454246
## 17	-0.2262961	-1.44939774	1.37565975	-0.28960865	-1.337784608
## 18	-0.1127499	0.39407030	0.38836278	-3.97985093	-0.410914404
## 19	2.9195668	1.58646124	-0.97612613	-0.78629766	-1.356288600
## 20	2.2998485	1.73396487	2.82423222	0.23281758	0.653038858
## 21	1.1501667	-0.13531015	-0.28506743	2.19770548	-0.084621572
## 22	-5.6594827	1.09730404	-0.10043541	0.05245484	0.689327990
## 23	-0.1011749	0.57911362	-0.71128354	0.44394773	-0.689939865
## 24	1.3836281	-1.95052341	2.98485490	0.35942784	0.744371276
## 25	0.2727756	-2.63013778	-1.83189535	-0.05207518	-0.803692524
## 26	4.0565577	-1.17534729	0.81690756	-1.66990720	2.895110075
## 27	0.8929694	-0.79236692	-1.26822542	0.57575615	-1.830793964
## 28	0.1514495	-1.44873320	-0.10857670	0.51040146	1.023229895
## 29	3.5592481	4.76202163	-0.75080576	-0.64692974	-0.309946510
## 30	-4.1184576	0.38073981	-1.43463965	-0.63330834	0.254715638
## 31	-0.6811731	-1.66926027	2.88645794	1.30977099	0.470913997
## 32	1.7157269	1.30836339	0.55971313	0.70557980	-0.331277622
## 33	-1.8860627	-0.59058174	-1.43570145	-0.18239089	-0.291863659
## 34	1.9526349	-0.52395429	0.75642216	-0.44289927	-0.723474420
## 35	1.5888864	3.12998571	1.73107199	1.68604766	-0.665406182
## 36	1.0709414	1.65628271	-0.79436888	1.85172698	-0.020031154
## 37	-4.1101715	-0.15766712	-2.36296974	0.56868399	2.469679496
## 38	-0.7254706	-2.89263339	0.36348376	0.50612576	-0.028157162
## 39	-3.3451254	0.95045293	-0.19551398	0.27716645	-0.487259213
## 40	-1.0644466	1.05265304	-0.82886286	0.12042931	0.645884788
## 41	1.4933989	-1.86712106	-1.81853582	1.06112429	-0.009855774
## 42	-0.6789284	-1.83156328	1.65435992	-0.95121379	-2.115630145
## 43	-2.4164258	0.46701087	-1.42808323	-0.41149015	0.867397522
## 44	2.2978729	-0.41865689	0.64422929	0.63462770	0.703116983
## 45	-2.9245282	1.19488555	3.35139309	1.48966984	-0.806659622
## 46	1.7654525	-0.95655926	-0.98576138	-1.05683769	-0.542466034
## 47	2.3125056	-2.56161119	1.58223354	-0.59863946	1.140712406


```

#Combining PCAs with crime
new.crime2 <- cbind(pca.chosen2, crime[,16]) %>%
  data.frame()
colnames(new.crime2)[6] <- "Crime"
new.model2 <- lm ( Crime ~ PC1 + PC2 + PC3 + PC4 + PC5 ,data = new.crime2 )
summary(new.model2)

##
## Call:
## lm(formula = Crime ~ PC1 + PC2 + PC3 + PC4 + PC5, data = new.crime2)
##
## Residuals:
##      Min       1Q   Median       3Q      Max
## -420.79 -185.01   12.21  146.24  447.86
##
## Coefficients:
##              Estimate Std. Error t value Pr(>|t|)
## (Intercept)   905.09      35.59   25.428 < 2e-16 ***
## PC1           65.22      14.67    4.447 6.51e-05 ***
## PC2           70.08      21.49    3.261 0.00224 **
## PC3          -25.19      25.41   -0.992 0.32725
## PC4          -69.45      33.37   -2.081 0.04374 *
## PC5          229.04      36.75    6.232 2.02e-07 ***
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
##
## Residual standard error: 244 on 41 degrees of freedom
## Multiple R-squared:  0.6452, Adjusted R-squared:  0.6019
## F-statistic: 14.91 on 5 and 41 DF,  p-value: 2.446e-08

```

Question 10.1

Using the same crime data set `uscrime.txt` as in Questions 8.2 and 9.1, find the best model you can using (a) a regression tree model, and (b) a random forest model.

```

#Question 10.1
library(rpart)
library(rpart.plot)
library(randomForest)
library(caret)
library(modelr)

#CART
# set up train and testing split
train <- createDataPartition(crime$Crime, p = .85, list = F)
# set up test and train datasets
crime.train <- crime[train,]
crime.test <- crime[-train,]
# check splits
dim(crime.train); dim(crime.test)

## [1] 43 16
## [1]  4 16

```

```

#Regression Tree
crime.tree <- train(
  Crime ~ .,
  data = crime.train,
  method = 'rpart',
  trControl = trainControl(method = 'boot_all', number = 10),
  metric = 'RMSE'
)

## Warning in nominalTrainWorkflow(x = x, y = y, wts = weights, info = trainInfo, :
## There were missing values in the apparent performance measures.


## Warning in nominalTrainWorkflow(x = x, y = y, wts = weights, info = trainInfo, :
## There were missing values in resampled performance measures.

crime.tree$finalModel

## n= 43
##
## node), split, n, deviance, yval
##      * denotes terminal node
##
## 1) root 43 6601051.0  910.2791
##    2) Po1< 10 32 1534630.0  747.8438
##      4) Po1< 7.15 19  654085.2  637.2105 *
##      5) Po1>=7.15 13  308103.2  909.5385 *
##    3) Po1>=10 11 1765868.0 1382.8180 *

prp(crime.tree$finalModel)

```



```

#Random Forest
crime.forest <- train(
  Crime ~ .,
  data = crime.train,
  method = 'rf',
  trControl = trainControl(method = 'boot_all', number = 10),
  metric = 'RMSE')
crime.forest$finalModel

```

```

##
## Call:
## randomForest(x = x, y = y, mtry = param$mtry)
##           Type of random forest: regression
##           Number of trees: 500
## No. of variables tried at each split: 8
##
##           Mean of squared residuals: 82700.7
##           % Var explained: 46.13

```

```

#Testing
crime.res1 <- crime.test %>%
  add_predictions(., crime.tree) %>%
  select('observations' = Crime, pred) %>%
  as.data.frame()
crime.res2 <- crime.test %>%
  add_predictions(., crime.forest) %>%
  select('observations' = Crime, pred) %>%

```

```

as.data.frame()

crime.res1

##      observations      pred
## 19             750 1382.8182
## 28             1216  909.5385
## 34              923  909.5385
## 46              508 1382.8182

crime.res2

##      observations      pred
## 19             750 1245.2786
## 28             1216  993.0923
## 34              923 1021.4021
## 46              508 1131.4758

postResample(obs = crime.res1$observations, pred = crime.res1$pred)

##      RMSE      Rsquared      MAE
## 561.2186717  0.7287986 456.8898601

postResample(obs = crime.res2$observations, pred = crime.res2$pred)

##      RMSE      Rsquared      MAE
## 416.3513502  0.4659483 360.0160583

```

Question 10.2

Describe a situation or problem from your job, everyday life, current events, etc., for which a logistic regression model would be appropriate. List some (up to 5) predictors that you might use.

I would consider logistic regression useful in predicting customer behavior on EC sites. The results would be buy (0) and don't buy(1). As for predictors, age, occupation, time spent on site would be considered good predictors.

Question 10.3

Using the GermanCredit data set `germancredit.txt`, use logistic regression to find a good predictive model for whether credit applicants are good credit risks or not.

```

set.seed(101)
credit <- read.table("germancredit.txt", header = FALSE)
str(credit)

## 'data.frame':   1000 obs. of  21 variables:
## $ V1 : chr  "A11" "A12" "A14" "A11" ...
## $ V2 : int   6 48 12 42 24 36 24 36 12 30 ...
## $ V3 : chr  "A34" "A32" "A34" "A32" ...
## $ V4 : chr  "A43" "A43" "A46" "A42" ...
## $ V5 : int  1169 5951 2096 7882 4870 9055 2835 6948 3059 5234 ...
## $ V6 : chr  "A65" "A61" "A61" "A61" ...
## $ V7 : chr  "A75" "A73" "A74" "A74" ...
## $ V8 : int   4 2 2 2 3 2 3 2 2 4 ...
## $ V9 : chr  "A93" "A92" "A93" "A93" ...
## $ V10: chr  "A101" "A101" "A101" "A103" ...

```

```
## $ V11: int 4 2 3 4 4 4 2 4 2 ...
## $ V12: chr "A121" "A121" "A121" "A122" ...
## $ V13: int 67 22 49 45 53 35 53 35 61 28 ...
## $ V14: chr "A143" "A143" "A143" "A143" ...
## $ V15: chr "A152" "A152" "A152" "A153" ...
## $ V16: int 2 1 1 1 2 1 1 1 1 2 ...
## $ V17: chr "A173" "A173" "A172" "A173" ...
## $ V18: int 1 1 2 2 2 2 1 1 1 1 ...
## $ V19: chr "A192" "A191" "A191" "A191" ...
## $ V20: chr "A201" "A201" "A201" "A201" ...
## $ V21: int 1 2 1 1 2 1 1 1 1 2 ...
```

```
credit$V21[credit$V21==1] <- 0
credit$V21[credit$V21==2] <- 1
```

#Dividing data

```
credit.part <- createDataPartition(credit$V21, times = 1, p = 0.7, list=FALSE)
head(credit.part)
```

```
##      Resample1
## [1,]         1
## [2,]         2
## [3,]         3
## [4,]         4
## [5,]         5
## [6,]         6
```

```
credit.train <- credit[credit.part,]
credit.valid <- credit[-credit.part,]
```

#model

```
credit.log <- glm(V21 ~ ., data = credit.train, family=binomial(link="logit"))
summary(credit.log)
```

```
##
## Call:
## glm(formula = V21 ~ ., family = binomial(link = "logit"), data = credit.train)
##
## Deviance Residuals:
##      Min       1Q   Median       3Q      Max
## -2.6141  -0.6484  -0.3327   0.5806   2.5401
##
## Coefficients:
##              Estimate Std. Error z value Pr(>|z|)
## (Intercept) -3.569e-01  1.342e+00  -0.266  0.790318
## V1A12        -2.958e-01  2.809e-01  -1.053  0.292252
## V1A13        -5.582e-01  4.807e-01  -1.161  0.245532
## V1A14        -1.507e+00  2.820e-01  -5.346   9e-08 ***
## V2           4.587e-02  1.215e-02   3.774  0.000160 ***
## V3A31         8.164e-02  6.777e-01   0.120  0.904113
## V3A32        -3.752e-01  5.371e-01  -0.699  0.484776
## V3A33        -1.110e+00  5.813e-01  -1.909  0.056260 .
## V3A34        -1.857e+00  5.579e-01  -3.329  0.000872 ***
## V4A41        -1.848e+00  4.817e-01  -3.836  0.000125 ***
## V4A410       -2.417e+00  9.679e-01  -2.497  0.012511 *
## V4A42        -9.125e-01  3.323e-01  -2.746  0.006029 **
```

```
## V4A43      -1.039e+00  3.116e-01  -3.334  0.000855 ***
## V4A44      -8.163e-01  9.009e-01  -0.906  0.364883
## V4A45      -3.390e-01  6.401e-01  -0.530  0.596388
## V4A46       1.099e-01  4.737e-01   0.232  0.816452
## V4A48      -1.674e+00  1.210e+00  -1.384  0.166426
## V4A49      -1.163e+00  4.208e-01  -2.764  0.005718 **
## V5         1.454e-04  5.535e-05   2.628  0.008594 **
## V6A62      -1.776e-01  3.428e-01  -0.518  0.604336
## V6A63       4.524e-01  4.577e-01   0.988  0.322956
## V6A64      -9.605e-01  6.456e-01  -1.488  0.136846
## V6A65      -8.524e-01  3.235e-01  -2.635  0.008419 **
## V7A72      -4.012e-01  5.588e-01  -0.718  0.472762
## V7A73      -6.514e-01  5.339e-01  -1.220  0.222457
## V7A74      -1.503e+00  5.815e-01  -2.585  0.009750 **
## V7A75      -8.686e-01  5.445e-01  -1.595  0.110686
## V8         3.831e-01  1.136e-01   3.372  0.000747 ***
## V9A92      -1.445e-01  5.564e-01  -0.260  0.795099
## V9A93      -6.374e-01  5.467e-01  -1.166  0.243694
## V9A94      -4.908e-01  6.419e-01  -0.765  0.444459
## V10A102     3.868e-01  4.966e-01   0.779  0.436073
## V10A103    -5.953e-01  4.893e-01  -1.217  0.223718
## V11        -2.731e-02  1.083e-01  -0.252  0.800898
## V12A122     5.001e-01  3.131e-01   1.597  0.110233
## V12A123     2.743e-01  2.899e-01   0.946  0.344022
## V12A124     6.561e-01  5.220e-01   1.257  0.208747
## V13        -1.923e-02  1.155e-02  -1.665  0.095899 .
## V14A142    -4.130e-01  5.260e-01  -0.785  0.432343
## V14A143    -8.298e-01  2.938e-01  -2.824  0.004744 **
## V15A152    -3.803e-01  2.869e-01  -1.326  0.184883
## V15A153    -6.739e-01  5.884e-01  -1.145  0.252124
## V16        7.446e-01  2.530e-01   2.944  0.003245 **
## V17A172     7.947e-01  8.275e-01   0.960  0.336845
## V17A173     5.865e-01  7.953e-01   0.737  0.460850
## V17A174     3.465e-01  8.183e-01   0.423  0.671940
## V18        1.495e-01  3.242e-01   0.461  0.644658
## V19A192    -3.794e-02  2.494e-01  -0.152  0.879062
## V20A202    -1.223e+00  9.110e-01  -1.343  0.179377
## ---
```

```
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
##
```

```
## (Dispersion parameter for binomial family taken to be 1)
##
```

```
## Null deviance: 848.32 on 699 degrees of freedom
## Residual deviance: 592.81 on 651 degrees of freedom
## AIC: 690.81
##
```

```
## Number of Fisher Scoring iterations: 5
```

```
#Confusion Matrix
```

```
creditPredict <- predict(credit.log, newdata=credit.valid[, -21], type="response")
Confusion.mat <- table(credit.valid$V21, round(creditPredict))
```

We can see that although sensitivity is quite high, Specifity isn't

```
Sensitivity <- Confusion.mat[1,1]/sum(Confusion.mat[1,])  
Sensitivity
```

```
## [1] 0.8786408
```

```
Specfityty <-Confusion.mat[2,2]/sum(Confusion.mat[2,])  
Specfityty
```

```
## [1] 0.4787234
```

Then we change the threshold

```
#setting second thresh hold  
threshold <- 0.7  
thres <- as.matrix(table(round(creditPredict > threshold), credit.valid$V21))  
names(dimnames(thres)) <- c("Predicted", "Observed")  
thres
```

```
##           Observed  
## Predicted    0    1  
##           0 197   76  
##           1   9   18
```

Below are the results for a different threshold, there is an obvious rise in specifity but also a slight loss in sensitivity.

```
Sensitivity2 <- thres[1,1]/sum(thres[1,])  
Sensitivity2
```

```
## [1] 0.7216117
```

```
Specfityty2 <-thres[2,2]/sum(thres[2,])  
Specfityty2
```

```
## [1] 0.6666667
```